

- [BFS Application](#)
  - [Requirements](#)
  - [Installation, Setup and Execution of the BFS code](#)
    - [Downloading](#)
    - [Setting up the environment](#)
    - [Building](#)
      - [Makefile Configuration](#)
    - [Test Scenarios](#)
  - [Runing tests](#)
- [Automation Scripts](#)
  - [List of scripts](#)
  - [Installing scripts](#)
  - [Script r-ify.sh](#)
    - [Description](#)
    - [Execution](#)
  - [Script r-compare.sh](#)
    - [Description](#)
    - [Execution](#)
  - [Script check-all.sh](#)
    - [Description](#)
    - [Execution](#)
  - [Script scorep\\_install.sh](#)
    - [Description](#)
    - [Execution](#)
- [Other](#)
  - [Profiling and Tracing](#)
  - [Current Limitations](#)
    - [Out-Of-Memory errors and CUDA memory size limitations:](#)
    - [Validation errors for the BFS runs](#)
  - [Troubleshooting](#)

- [About this document](#)
- [License](#)
- [External Resources](#)
- 

# BFS Application

---

## Requirements

---

- Compiler with std c++11 support (e.g GNU C/C++ v4.8.1+)
- UNIX-like OS with CUDA 6+ support

## Installation, Setup and Execution of the BFS code

---

## Downloading

Create an account in <https://bitbucket.org> Request access for bfs\_multinode repository

Available repositories are: (Replace the token Your\_Bitbucket\_User)

- [\(Updated daily\)](https://Your_Bitbucket_User@bitbucket.org/julianromera/bfs_multinode.git)
- [\(Updated weekly\)](https://Your_Bitbucket_User@bitbucket.org/jyoung3131/bfs_multinode.git)

```
$ REPLACE the token Your_Bitbucket_User  
$ git clone https://Your_Bitbucket_User@bitbucket.org/julianromera/bfs_multinode.git  
$ git checkout -b architectural_tuning  
$ git pull origin architectural_tuning
```

## Setting up the environment

1. Add CUDA libraries and binaries to your path:

Add this text to your ~/.bashrc

```
UPDATE NEXT LINE  
export CUDA_PATH=/usr/local/cuda-7.0  
export PATH=$CUDA_PATH/bin:$PATH  
export LD_LIBRARY_PATH=$CUDA_PATH/lib64:$CUDA_PATH/lib:$CUDA_PATH/extras/CUPT
```

## 1. In case of using EXTOLL add these lines to your ~/.bashrc

```
Adjust next line to your actual path  
export EXTOLL_HOME=/extoll2  
if [ -d /extoll2 ]; then  
    source $EXTOLL_HOME/extoll_env.bash  
    # ensure this path is correct  
    export MPI_PATH=$EXTOLL_HOME/mpi/openmpi-1.6.4  
    export PATH=$MPI_PATH/bin:$PATH  
  
    export LD_LIBRARY_PATH=$EXTOLL_HOME/lib:$MPI_PATH/lib:$LD_LIBRARY_PATH  
    ulimit -l 4194304  
    ulimit -s unlimited  
fi
```

## Building

The code to compile is in the folder `cpu_2d/`. It is built using `Make`.

The currently available Makefiles are:

- `Makefile.gcc`
- `Makefile.gcc.keeneland`

The Makefiles in the list below are currently being developed:

- `Makefile.Intel`
- `Makefile.cpu`

The file `Makefile` is a symbolic link to the `Makefile.____` being used. It may be created/edited using:

```
$ cd bfs_multinode/cpu_2d  
$ rm -f Makefile  
$ ln -s Makefile.gcc Makefile
```

```
$ cd bfs_multinode/cpu_2d  
$ make  
$ cd ../eval
```

## Makefile Configuration

The Makefiles in the `cpu_2d/` folder are configurable. Edit the Makefile file to access the documentation for the variables. The documentation is placed in the header of the file.

The default variables/ values for Makefile.gcc.Keeneland are:

```
nvidia_architecture      := "auto"  
nvidia_ptxas_optimize    := "no"  
manual_profiler_cuda     := "no"  
manual_profiler_other_compilers := "yes"  
openmp_on_cuda           := "no"  
openmp_on_other_compilers := "no"  
custom_openmpi            := "no"  
custom_openmpi_basedir   := /home/CHANGE_ME/openmpi  
scorep_profiler_enable    := "no"  
scorep_profiler_automatic_instrumentation := "yes"  
scorep_custom              := "yes"  
scorep_custom_basedir     := /home/CHANGE_ME/scorep  
use_avx_instructions     := "yes"  
enable_simd_compression    := "yes"  
enable_simd_compression_benchmark := "no"  
debug                     := "no"  
quiet_output               := "no"  
use_cuda                  := "yes"  
code_optimization_level    := "O4"  
code_optimization_flags    := -funroll-loops -flio
```

## Test Scenarios

Test scenarios are in the folder `eval/`

SBatch relevant Test-cases in the previous list are:

- o4p2n-roptim.rsh [SCALE\_FACTOR]
- o4p2n-coptim.rsh [SCALE\_FACTOR]
- o4p2n-noptim.rsh [SCALE\_FACTOR]
- o9p8n.rsh [SCALE\_FACTOR]

- o16p8n.rsh [SCALE\_FACTOR]

The name in these files explains what they do. For example:

The part "4p2n" indicates 4 Processes will be run in 2 Nodes. These scripts are run through Slurm. "16p8n" would indicate 16 processes distributed in 8 Nodes.

## Running tests

---

Ensure that the Slurm queue is empty

```
$ squeue  
JOBID PARTITION NAME USER ST TIME NODES NODELIST(REASON)
```

Use Sbatch to execute the .rsh file. Add a numeric argument at the end (Sf)

```
$ sbatch o4p2n_roptim.rsh 21  
Submitted batch job 423  
$ sbatch o4p2n_coptim.rsh 21vv  
Submitted batch job 424  
$ sbatch o4p2n_noptim.rsh 21  
Submitted batch job 425
```

## Automation Scripts

---

### List of scripts

---

- r-ify.sh
- r-compare.sh
- check-all.sh
- scorep\_install.sh

### Installing scripts

---

Scripts are located under the `scripts/` folder

To install

- r-ify.sh
- r-compare.sh
- check-all.sh

Run:

```
$ cd bfs_multinode
$ # Make a symlink from eval to the script file
$ cd eval
$ ln -s ../scripts/r-ify.sh r-ify.sh
$ ln -s ../scripts/r-compare.sh r-compare.sh
$ ln -s ../scripts/check-all.sh check-all.sh
$ chmod u+x *.sh
```

To install (Useful for Profiling and Tracing)

- scorep\_install.sh

Run:

```
cp scripts/scorep_install.sh ~
cd ~
chmod u+x scorep_install.sh
./scorep_install.sh
...
```

## Script r-ify.sh

---

### Description

It uses the execution traces to generate R-code. This R-code (once run) shows the time measurements of the Phases of the BFS code. This script uses result files with same Scale Factor. The results are represented as a Barplot.

### Execution

```
$ ./r-ify.sh 423 424 425  
-> File slurm-423.out. Validation passed (Tasks: 4, GPUS/Task: 1, Scale Factor: 21).  
-> File slurm-424.out. Validation passed (Tasks: 4, GPUS/Task: 1, Scale Factor: 21).  
-> File slurm-425.out. Validation passed (Tasks: 4, GPUS/Task: 1, Scale Factor: 21).  
Enter new labels for the X-Axe? (y/n) [n] y  
Enter a total 3 label(s) between quotes. Separate them with spaces: "4p2n-Roptim" "4p2n-Cop  
-> Created file "file-423-424-425.r".  
-> R-Code successfully generated.
```

Open the file `file-423-424-425.r` with your R editor and run the code.

## Script r-compare.sh

---

### Description

As the previous script, this also uses the execution traces to generate R-code. This differs from the previous one in that it can compare several files from several Scale Factors. Results are visualized as a Lineplot.

### Execution

```
$ ./r-ify.sh  
(to be completed)
```

## Script check-all.sh

---

### Description

This script automatizes the execution of tests for different Scale Factors.

### Execution

```
$ check-all.sh 15 30
```

This will run the tests with format `o*.rsh` in the `eval/` folder for Scale Factors 15 to 30.

Process is shown in ncurses-like format.

## Script scorep\_install.sh

---

### Description

This script Download, Uncompress, Builds and Install locally Score-P, Scalasca and CUBE, and a compatible version of OpenMPI. This script is configurable. The configuration may be changed editing the script.

Root priviledges are not required to run the script.

### Execution

```
$ cd $HOME  
$ ./scorep_install.sh
```

## Other

---

### Profiling and Tracing

This BFS application allows the code to be instrumented in Zones using Score-P with very low overhead. This requires Score-P and Scalasca to be installed in the system. The results may be analyzed either visually (using CUBE) or through console using scorep-score

These tools may be installed locally (No privileged user is needed) using the scorep\_install.sh aforementioned.

The names of the instrumentable Zones are listed below. Others may be added if needed.

```
BFSRUN_region_vertexBroadcast  
BFSRUN_region_nodesTest  
BFSRUN_region_localExpansion  
BFSRUN_region_testSomethingHasBeenCalled  
BFSRUN_region_columnCommunication  
BFSRUN_region_rowCommunication
```

The first step is to update the system variables. This may be done either on .bashrc or in a separate script.

Update the paths in the variables below

```
$ cat >> ~/.bashrc << EOF  
export G500_ENABLE_RUNTIME_SCALASCA=yes  
  
export SCOREP_CUDA_BUFFER=48M  
export SCOREP_CUDA_ENABLE=no  
export SCOREP_ENABLE_PROFILING=true  
export SCOREP_ENABLE_TRACING=false  
export SCOREP_PROFILING_FORMAT=CUBE4  
export SCOREP_TOTAL_MEMORY=12M  
export SCOREP_VERBOSE=no  
export SCOREP_PROFILING_MAX_CALLPATH_DEPTH=330  
  
export LD_LIBRARY_PATH=$HOME/cube/lib:$LD_LIBRARY_PATH  
export PATH=$HOME/cube/bin:$PATH  
  
export LD_LIBRARY_PATH=$HOME/scorep/lib:$LD_LIBRARY_PATH  
export PATH=$HOME/scorep/bin:$PATH  
  
export LD_LIBRARY_PATH=$HOME/scalasca/lib:$LD_LIBRARY_PATH  
export PATH=$HOME/scalasca/bin:$PATH  
  
export MPI_PATH=/home/jromera/openmpi  
export PATH=$MPI_PATH/bin:$CUDA_PATH/bin:$PATH  
  
export LD_LIBRARY_PATH=$MPI_PATH/lib:$CUDA_PATH/lib64:$CUDA_PATH/lib64/stubs:$  
export LD_LIBRARY_PATH=$HOME/scorep/lib:$LD_LIBRARY_PATH  
export PATH=$HOME/scorep/bin:$PATH  
EOF
```

The variable `G500_ENABLE_RUNTIME_SCALASCA` set to yes will enable the required runtime instrumentor of Scalasca.

Results will be stored on a folder with format `scorep-*` in the `/eval` folder.

To instrument graphically with CUBE run:

```
$ cd eval(scorep-____FOLDER_NAME____  
$ cube profile.cubex
```

To instrument through the console run:

```
$ cd eval(scorep-____FOLDER_NAME____  
$ scorep-score -r profile.cubex
```

## Current Limitations

---

### Out-Of-Memory errors and CUDA memory size limitations:

For some high Score-Factors (e.g 22 as of the day of writing this guide), the resulting Slurm trace will be:

Using SCALE-FACTOR 22

Tue Jul 14 15:42:51 CEST 2015

===== JOB MAP =====

Data for node: creek01 Num procs: 2

Process OMPI jobid: [1404,1] Process rank: 0

Process OMPI jobid: [1404,1] Process rank: 2

Data for node: creek02 Num procs: 2

Process OMPI jobid: [1404,1] Process rank: 1

Process OMPI jobid: [1404,1] Process rank: 3

=====  
row slices: 2, column slices: 2

graph\_generation: 7.749108 s

Input list of edges generated.

6.710886e+07 edge(s) generated in 8.499692s (7.895447 Medges/s on 4 processor(s))

Adjacency Matrix setup.

2.956432e+06 edge(s) removed, because they are duplicates or self loops.

1.283049e+08 unique edge(s) processed in 18.308235s (7.008041 Medges/s on 4 processor(s)  
[./b40c/graph/bfs/csr\_problem\_2d.cuh, 697] CsrProblem cudaMalloc frontier\_queues.d\_values  
[cuda/cuda\_bfs.cu, 486] Reset error. (CUDA error 2: out of memory)

MPI\_ABORT was invoked on rank 0 in communicator MPI\_COMM\_WORLD  
with errorcode 1.

## Validation errors for the BFS runs

Looking at your resulting Sbatch trace you may see:

"Validation: failed" / or an output trace similar to the one below:

The issue seems to be related with compiler version errors. The version of the compiler used for the CUDA code may not be met. See the Requirements section above.

```
$ cat slurm-JOBID.out
...
BFS Iteration 24: Finished in 0.111173s
max. local exp.: 0.004383s(3.942371%)
max. queue handling: 0.056909s(51.189917%)
est. rest: 0.049881s(44.867712%)
max. row com.: 0.038369s(34.513278%)
max. col com.: 0.069516s(62.529783%)
max. pred. list. red: 0.007413s(6.667911%)
(3:0) Edge [26849(2273):244(244)] with wrong levels: 3 -1
Validation of iteration 24 finished in 0.022564s
Result: Invalid 524279 Edge(s) processed, 4.715888 MTeps
(0:3) Edge [244(244):26849(2273)] with wrong levels: -1 3
(0:3) Edge [244(244):29960(5384)] with wrong levels: -1 4
(3:0) Edge [29960(5384):244(244)] with wrong levels: 4 -1
BFS Iteration 25: Finished in 0.096601s
max. local exp.: 0.000905s(0.937128%)
max. queue handling: 0.048780s(50.496083%)
est. rest: 0.046916s(48.566789%)
max. row com.: 0.032854s(34.009833%)
max. col com.: 0.054057s(55.958921%)
max. pred. list. red: 0.009008s(9.324883%)
Validation of iteration 25 finished in 0.020493s
Result: Valid 524280 Edge(s) processed, 5.427272 MTeps
invalid_level
invalid_level
...
```

## Troubleshooting

- Problem: In the .out file of Slurm/ Sbatch execution I get the text:

```
S=C=A=N: Abort: No SCOREP instrumentation found in target ..../cpu_2d/g500
```

- Solution:

The instrumentation is activated for the runtime execution (i.e: the binary is being run prefixed with scalasca).

Disable it with:

```
$ export G500_ENABLE_RUNTIME_SCALASCA=no
```

# About this document

---

- Version 1.0
- Last revision: 12th August 2015

## License

---

This code contains a subset of Duane Merrill's BC40 repository of GPU-related functions, including his BFS implementation used in the paper, Scalable Graph Traversals.

All copyrights reserved to their original owners.

## External Resources

---

<https://www.rstudio.com/products/RStudio/>